

▼ Installing packages

```
!pip install langchain_google_genai
```

```
Requirement already satisfied: langchain_google_genai in /usr/local/lib/python3.12/dist-packages (4.2.0)
Requirement already satisfied: filetype<2.0.0,>=1.2.0 in /usr/local/lib/python3.12/dist-packages (from langchain_google_genai) (1.2.0)
Requirement already satisfied: google-genai<2.0.0,>=1.56.0 in /usr/local/lib/python3.12/dist-packages (from langchain_google_genai) (1.56.0)
Requirement already satisfied: langchain-core<2.0.0,>=1.2.5 in /usr/local/lib/python3.12/dist-packages (from langchain_google_genai) (1.2.5)
Requirement already satisfied: pydantic<3.0.0,>=2.0.0 in /usr/local/lib/python3.12/dist-packages (from langchain_google_genai) (2.12.3)
Requirement already satisfied: aiohttp<5.0.0,>=4.8.0 in /usr/local/lib/python3.12/dist-packages (from google-genai<2.0.0,>=1.56.0->langchain)
Requirement already satisfied: google-auth<3.0.0,>=2.47.0 in /usr/local/lib/python3.12/dist-packages (from google-auth[requests]<3.0.0,>=2.47.0->langchain)
Requirement already satisfied: requests<3.0.0,>=2.28.1 in /usr/local/lib/python3.12/dist-packages (from google-genai<2.0.0,>=1.56.0->langchain)
Requirement already satisfied: tenacity<9.2.0,>=8.2.3 in /usr/local/lib/python3.12/dist-packages (from google-genai<2.0.0,>=1.56.0->langchain)
Requirement already satisfied: websockets<15.1.0,>=13.0.0 in /usr/local/lib/python3.12/dist-packages (from google-genai<2.0.0,>=1.56.0->langchain)
Requirement already satisfied: typing_extensions<5.0.0,>=4.11.0 in /usr/local/lib/python3.12/dist-packages (from google-genai<2.0.0,>=1.56.0->langchain)
Requirement already satisfied: distro<2,>=1.7.0 in /usr/local/lib/python3.12/dist-packages (from google-genai<2.0.0,>=1.56.0->langchain)
Requirement already satisfied: sniffio in /usr/local/lib/python3.12/dist-packages (from google-genai<2.0.0,>=1.56.0->langchain)
Requirement already satisfied: jsonpatch<2.0.0,>=1.33.0 in /usr/local/lib/python3.12/dist-packages (from langchain-core<2.0.0,>=1.2.5->langchain)
Requirement already satisfied: langsmith<1.0.0,>=0.3.45 in /usr/local/lib/python3.12/dist-packages (from langchain-core<2.0.0,>=1.2.5->langchain)
Requirement already satisfied: packaging<26.0.0,>=23.2.0 in /usr/local/lib/python3.12/dist-packages (from langchain-core<2.0.0,>=1.2.5->langchain)
Requirement already satisfied: pyyaml<7.0.0,>=5.3.0 in /usr/local/lib/python3.12/dist-packages (from langchain-core<2.0.0,>=1.2.5->langchain)
Requirement already satisfied: uuid-utils<1.0.0,>=0.12.0 in /usr/local/lib/python3.12/dist-packages (from langchain-core<2.0.0,>=1.2.5->langchain)
Requirement already satisfied: annotated-types<=0.6.0 in /usr/local/lib/python3.12/dist-packages (from pydantic<3.0.0,>=2.0.0->langchain)
Requirement already satisfied: pydantic-core==2.41.4 in /usr/local/lib/python3.12/dist-packages (from pydantic<3.0.0,>=2.0.0->langchain)
Requirement already satisfied: typing-inspection<=0.4.2 in /usr/local/lib/python3.12/dist-packages (from pydantic<3.0.0,>=2.0.0->langchain)
Requirement already satisfied: idna>=2.8 in /usr/local/lib/python3.12/dist-packages (from aiohttp<5.0.0,>=4.8.0->google-genai<2.0.0,>=1.56.0->langchain)
Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.12/dist-packages (from google-auth<3.0.0,>=2.47.0->google-genai)
Requirement already satisfied: cryptography>=38.0.3 in /usr/local/lib/python3.12/dist-packages (from google-auth<3.0.0,>=2.47.0->google-genai)
Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.12/dist-packages (from google-auth<3.0.0,>=2.47.0->google-auth[re])
Requirement already satisfied: certifi in /usr/local/lib/python3.12/dist-packages (from https://<1.0.0,>=0.28.1->google-genai<2.0.0,>=1.56.0->langchain)
Requirement already satisfied: httpcore==1.* in /usr/local/lib/python3.12/dist-packages (from https://<1.0.0,>=0.28.1->google-genai<2.0.0,>=1.56.0->langchain)
Requirement already satisfied: h11>=0.16 in /usr/local/lib/python3.12/dist-packages (from httpcore==1.*->https://<1.0.0,>=0.28.1->google-genai)
Requirement already satisfied: jsonpointer>=1.9 in /usr/local/lib/python3.12/dist-packages (from jsonpatch<2.0.0,>=1.33.0->langchain)
Requirement already satisfied: orjson>=3.9.14 in /usr/local/lib/python3.12/dist-packages (from langsmith<1.0.0,>=0.3.45->langchain)
Requirement already satisfied: requests-toolbelt>=1.0.0 in /usr/local/lib/python3.12/dist-packages (from langsmith<1.0.0,>=0.3.45->langchain)
Requirement already satisfied: zstandard>=0.23.0 in /usr/local/lib/python3.12/dist-packages (from langsmith<1.0.0,>=0.3.45->langchain)
Requirement already satisfied: charset_normalizer<4,>=2 in /usr/local/lib/python3.12/dist-packages (from requests<3.0.0,>=2.28.1->google-genai)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.12/dist-packages (from requests<3.0.0,>=2.28.1->google-genai)
Requirement already satisfied: cffi>=1.12 in /usr/local/lib/python3.12/dist-packages (from cryptography>=38.0.3->google-auth<3.0.0,>=2.4)
Requirement already satisfied: pyasn1<0.7.0,>=0.6.1 in /usr/local/lib/python3.12/dist-packages (from pyasn1-modules>=0.2.1->google-auth[re])
Requirement already satisfied: pyparser in /usr/local/lib/python3.12/dist-packages (from cffi>=1.12->cryptography>=38.0.3->google-auth[re])
```

▼ Setup your API key

To run the following cell, your API key must be stored it in a Colab Secret named VERTEX_API_KEY.

1. Look for the key icon on the left panel of your colab.
2. Under Name, create VERTEX_API_KEY.
3. Copy your key to Value.

```
from google.colab import userdata
GEMINI_VERTEX_API_KEY = userdata.get('VERTEX_API_KEY')
```

▼ Downloading receipts.zip

The codes below download and unzip receipts.zip from Google Drive. receipts.zip contains all images from the Fusion folder on BlackBoard.

```
import gdown
file_id = "1oe2FZd3ZT07nrDqjCafNvxic108oF8JF"
download_url = f"https://drive.google.com/uc?id={file_id}"
gdown.download(download_url, "receipts.zip", quiet=False)
```

```
Downloading...
From: https://drive.google.com/uc?id=1oe2FZd3ZT07nrDqjCafNvxic108oF8JF
To: /content/receipts.zip
100%|██████████| 1.61M/1.61M [00:00<00:00, 148MB/s]
'receipts.zip'
```

```
!unzip receipts.zip
```

```
Archive: receipts.zip
replace receipt1.jpg? [y]es, [n]o, [A]ll, [N]one, [r]ename:
```

▼ 1. Helper functions

We need two functions

- `image_to_base64` convert your jpg image into Base64 encoded string (basically a sequence of 64 characters to make your image easily transferred via API)
- `get_image_data_url` takes your jpg image, converting them into base64 string and construct a suitable input for GEMINI api call.

```
import base64
import mimetypes

# Helper function to read and encode image
def image_to_base64(img_path):
    with open(img_path, "rb") as img_file:
        return base64.b64encode(img_file.read()).decode('utf-8')

# Helper function to encode local file to Base64 Data URL
def get_image_data_url(image_path):
    # Guess the mime type (e.g., image/png, image/jpeg) based on file extension
    mime_type, _ = mimetypes.guess_type(image_path)
    if mime_type is None:
        mime_type = "image/png" # Default fallback

    encoded_string = image_to_base64(image_path)

    # Construct the Data URL
    return f"data:{mime_type};base64,{encoded_string}"
```

```
from langchain_google_genai import ChatGoogleGenerativeAI
llm = ChatGoogleGenerativeAI(
    model="gemini-2.5-flash",
    api_key=userdata.get('VERTEX_API_KEY'), # Ensure this key is set in Colab secrets
    temperature=0,
    vertexai=True
)
```

Display jpg images. Alternatively, open the folder icon on the left pannel to see the images.

```
from IPython.display import HTML, display
import glob, os

image_paths = glob.glob("*.jpg")
image_paths.sort()
html_content = '<div style="display: flex; flex-wrap: wrap; gap: 20px;">'

for path in image_paths:
    b64 = image_to_base64(path)
    filename = os.path.basename(path) # Clean up path to show just the name

    # Create a vertical column for each image + text
    html_content += f'''
        <div style="display: flex; flex-direction: column; align-items: center;">
            
            <span style="font-family: monospace; font-size: 14px;">{filename}</span>
        </div>
    ,,
```

html_content += '</div>'
display(HTML(html_content))



receipt1.jpg



receipt2.jpg



receipt3.jpg



receipt4.jpg



receipt5.jpg



receipt6.jpg



receipt7.jpg

2. Image input to Gemini

Different from text, image needs to be converted into base64 encoded string and then formated into url before inputting to the language model. This is convenient for image-type input to be transferred through the API.

This project builds a multimodal LLM pipeline to extract totals from multiple receipt images and answer two queries: **Q1** (total paid after discounts) and **Q2** (total before discounts).

Each receipt is processed individually. The image is converted to a data URL, sent to the LLM with a strict output format (`FINAL=...` `ORIGINAL=...`), and the numbers are extracted by regex and summed into `query1_answer` and `query2_answer` (single floats).

If `ORIGINAL` is missing, the code asks the LLM for `DISCOUNT` and estimates the pre-discount total as `FINAL + DISCOUNT`. Results are stored and printed in a summary table.

A separate classifier chain outputs only `Q1`, `Q2`, or `REJECT`, so irrelevant user queries are rejected.

```
from google.colab import auth
auth.authenticate_user()

import os, re
from langchain_core.prompts import ChatPromptTemplate

# ----- Process receipts one by one to reduce errors -----
query1_answer = 0.0
query2_answer = 0.0

results = [] # store per-receipt outputs for summary

for i, image_path in enumerate(image_paths):
    image_url = get_image_data_url(image_path)

    # Create a per-receipt prompt
    per_receipt_prompt = ChatPromptTemplate.from_messages([
        ("system",
            "You are a cashier analyzing ONE receipt. Extract the\n"
            "1. Find the 'Total', 'Amount Payable', or 'Grand Total'\n"
            "2. Find the total BEFORE any discounts. If not shown,\n"
            "3. If you see discounts (negative numbers or 'Discount')\n"
            "Return ONLY in this format: FINAL=XX.XX, ORIGINAL=YY.\n"),
        ("user", ""),
        ("assistant", "")])
```

```

        {"type": "text", "text": "What is the final total and original total on this receipt?"},  

        {"type": "image_url", "image_url": {"url": image_url}}  

    ])  

])  
  

per_receipt_chain = per_receipt_prompt | llm  

response = per_receipt_chain.invoke({})  

text = str(response.content)  
  

# Extract numbers using regex (tolerant)  

final_match = re.search(r'FINAL\s*[:=]\s*\[$€£¥]\s*\s*([\d]+(?:\.\[\d]+)?', text, re.IGNORECASE)  

original_match = re.search(r'ORIGINAL\s*[:=]\s*\[$€£¥]\s*\s*([\d]+(?:\.\[\d]+)?', text, re.IGNORECASE)  
  

final_total = None  

original_total = None  

original_est = None  

discount = None  
  

if final_match:  

    final_total = float(final_match.group(1))  

    query1_answer += final_total  
  

if original_match:  

    original_total = float(original_match.group(1))  

    query2_answer += original_total  

else:  

    # If ORIGINAL is not found, ask the LLM to compute total discount and estimate ORIGINAL  

    if final_total is not None:  

        discount_prompt = ChatPromptTemplate.from_messages([
            ("system",
             "Find the total discount amount on this receipt.\n"
             "Return ONLY in this format: DISCOUNT=XX.XX"),
            ("human", [
                {"type": "text", "text": "What is the total discount amount on this receipt?"},  

                {"type": "image_url", "image_url": {"url": image_url}}
            ])
        ])
        discount_response = (discount_prompt | llm).invoke({})
        disc_text = str(discount_response.content)  
  

        disc_match = re.search(r'DISCOUNT\s*[:=]\s*\[$€£¥]\s*\s*([\d]+(?:\.\[\d]+)?', disc_text, re.IGNORECASE)
        if not disc_match:
            disc_match = re.search(r' ([\d]+(?:\.\[\d]+)?', disc_text) # fallback  
  

        if disc_match:
            discount = float(disc_match.group(1))
            original_est = final_total + discount
            query2_answer += original_est  
  

results.append({
    "No": i + 1,
    "Receipt": os.path.basename(image_path),
    "FINAL": final_total,
    "ORIGINAL": original_total,
    "ORIG_EST": original_est,
    "RAW": text
})  
  

# ======  

# Summary output (not split)  

# ======  

print("\n" + "="*80)
print("Summary (all receipts):")
print(f'{No} :<4} {Receipt:<25} {FINAL:>12} {ORIGINAL:>14} {ORIG_EST:>12}')
for r in results:
    f = "" if r["FINAL"] is None else f'{r["FINAL"]:.2f}'
    o = "" if r["ORIGINAL"] is None else f'{r["ORIGINAL"]:.2f}'
    e = "" if r["ORIG_EST"] is None else f'{r["ORIG_EST"]:.2f}'
    print(f'{r["No"]:<4} {r["Receipt"]:<25} {f:>12} {o:>14} {e:>12}'")  
  

print("-"*80)
print(f"Query 1 total spent (after discount): {query1_answer:.2f}")
print(f"Query 2 total spent (before discount): {query2_answer:.2f}")
print("="*80)  
  

# ======  

# Add-on: Reject irrelevant queries  

# ======  

def answer_user_query(user_query: str):
    cls_prompt = ChatPromptTemplate.from_messages([
        ("system",
         "Classify the user query. Output ONLY one token: Q1, Q2, or REJECT.\n"
         "Q1 = total spent in total AFTER discounts.\n")
    ])

```

```

        "Q2 = total that would have been paid WITHOUT discounts (before discounts).\n"
        "Anything else = REJECT.",
        ("human", "{q}")
    ])

tag = (cls_prompt | llm).invoke({"q": user_query}).content.strip().upper()

if "Q1" in tag:
    return query1_answer
elif "Q2" in tag:
    return query2_answer
else:
    return "REJECTED" # if needed: "NONE"

```

```
=====
Summary (all receipts):
No. Receipt      FINAL      ORIGINAL      ORIG_EST
1  receipt1.jpg   394.70     480.40
2  receipt2.jpg   316.10     392.20
3  receipt3.jpg   140.80     160.10
4  receipt4.jpg   514.00     590.80
5  receipt5.jpg   102.30     107.70
6  receipt6.jpg   190.80     221.20
7  receipt7.jpg   315.60     396.00
-----
Query 1 total spent (after discount): 1974.30
Query 2 total spent (before discount): 2348.40
=====
```

3. Evaluation Code

- Make sure your LLM return a single float as the answer, stored in `query1_answer` and `query2_answer`
- Run the following code blocks: (1) If the blocks does not return any error, then your chain design is correct. Otherwise, please check your chain design.
- Do not modify `query1_costs` and `query2_costs`

```

def test_query(answer, ground_truth_costs):
    # Convert string to float if necessary
    if isinstance(answer, str):
        answer = float(answer)

    # Calculate the ground truth sum once for clarity
    expected_total = sum(ground_truth_costs)

    # Check if the answer is within +/- $2 of the expected total
    assert abs(answer - expected_total) <= 2

```

Run the following code block to evaluate query 1:

How much did I spend in total for these bills?

```
query1_costs = [394.7, 316.1, 140.8, 514.0, 102.3, 190.8, 315.6] # do not modify this
test_query(query1_answer, query1_costs)
```

Run the following code block to evaluate query 2:

How much would I have had to pay without the discount?

```
query2_costs = [480.20, 392.20, 160.10, 590.80, 107.70, 221.20, 396.00] # do not modify this
test_query(query2_answer, query2_costs)
```

```
sum([480.20, 392.20, 160.10, 590.80, 107.70, 221.20, 396.00])
```

```
2348.2
```