

Homework 2

- ✓ Set up
- ✓ Installing packages

```
!pip install requests PyPDF2 gdown
!pip install 'markitdown[pdf]'
!pip install langchain_mcp_adapters langchain_google_genai langchain-openai
```

[显示隐藏的输出项](#)

- ✓ Setup your API key

To run the following cell, your API key must be stored it in a Colab Secret named `VERTEX_API_KEY`.

1. Look for the key icon on the left panel of your colab.
2. Under `Name`, create `VERTEX_API_KEY`.
3. Copy your key to `Value`.

If you cannot use `VERTEX_API_KEY`, you can use deepseek models via `DEEPEEK_API_KEY`. It does not affect your score.

```
from google.colab import userdata
DEEPEEK_API_KEY = userdata.get('DEEPEEK_API_KEY')
```

- ✓ Download sample CVs

- ✓ Downloading sample_cv.pdf

The codes below download the sample CV

```
import os
import gdown

folder_id = "1adYKq7gSSczFP3iikfA8Er-HSZP6W7D"
folder_url = f"https://drive.google.com/drive/folders/{folder_id}"

output_dir = "downloaded_cv"
os.makedirs(output_dir, exist_ok=True)

gdown.download_folder(
    url=folder_url,
    output=output_dir,
    quiet=False,
    use_cookies=False
)
```

[显示隐藏的输出项](#)

```
# =====
# Load and display all CV PDFs in order
# =====
import os
from markitdown import MarkItDown

cv_dir = "downloaded_cv"

# Initialize MarkItDown
md = MarkItDown(enable_plugins=False)

# Collect and sort PDFs numerically
pdf_files = sorted([
    f for f in os.listdir(cv_dir) if f.lower().endswith(".pdf"),
    key=lambda x: int("".join(filter(str.isdigit, x))) # CV_1.pdf → 1
])
```

```
all_csvs = []

for pdf_name in pdf_files:
    pdf_path = os.path.join(cv_dir, pdf_name)
    result = md.convert(pdf_path)

    all_csvs.append({
        "file": pdf_name,
        "text": result.text_content
    })

print("-" * 80)
print(f" {pdf_name}")
print("-" * 80)
print(result.text_content)
print("\n\n")
```

[显示隐藏的输出项](#)

- ✓ Connect to our MCP server

Documentation about MCP: <https://modelcontextprotocol.io/docs/getting-started/intro>.

Using MCP servers in Langchain <https://docs.langchain.com/oss/python/langchain/mcp>.

- ✓ Check which tools that the MCP server provide

```
import asyncio
import json
from langchain_mcp_adapters.client import MultiServerMCPClient

client = MultiServerMCPClient({
    "social_graph": {
        "transport": "http",
        "url": "https://ftec5660.ngrok.app/mcp",
        "headers": {""ngrok-skip-browser-warning": "true"}
    }
})

mcp_tools = await client.get_tools()
for tool in mcp_tools:
    print(tool.name)
    print(tool.description)
    print(tool.args)
    print("\n\n")
```

[显示隐藏的输出项](#)

- ✓ A simple agent using tools from the MCP server

```
from google.colab import userdata
from langchain_openai import ChatOpenAI

DEEPEEK_API_KEY = userdata.get('DEEPEEK_API_KEY')

llm = ChatOpenAI(
    model="deepseek-chat",
    api_key=DEEPEEK_API_KEY,
    base_url="https://api.deepseek.com/v1",
    temperature=0,
)
```

```
import json
import re

def _tool_to_obj(tool_result):
    if isinstance(tool_result, list) and tool_result and isinstance(tool_result[0], dict) and "text" in tool_result[0]:
        txt = tool_result[0]["text"]
        try:
            return json.loads(txt)
        except:
            return txt
    return tool_result
```

```

def _pick_best_candidate(items, name_key="name"):
    """
    Select the most similar profile by giving priority to exact matches. If multiple candidates remain, choose the
    if not items:
        return None
    return sorted(
        items,
        key=lambda x: (x.get("match_type") != "exact", -x.get("years_experience", 0))
    )[0]

async def fetch_linkedin_best(tools, name: str, location_hint: str | None):
    t_search = next(t for t in tools if t.name == "search_linkedin_people")
    t_get     = next(t for t in tools if t.name == "get_linkedin_profile")

    # 1) Start by searching with the location included
    rl = await t_search.ainvoke({
        "q": name,
        "location": location_hint,
        "industry": None,
        "limit": 10,
        "fuzzy": True
    })
    people = _tool_to_obj(rl)

    # 2) If no result is found, remove the location and search again so that we return the most similar match
    if not people:
        r2 = await t_search.ainvoke({
            "q": name,
            "location": None,
            "industry": None,
            "limit": 10,
            "fuzzy": True
        })
        people = _tool_to_obj(r2)

    if not people:
        return None

    best = _pick_best_candidate(people)
    prof_raw = await t_get.ainvoke({"person_id": best["id"]})
    profile = _tool_to_obj(prof_raw)
    return {"candidate": best, "profile": profile}

async def fetch_facebook_best(tools, name: str):
    t_search = next(t for t in tools if t.name == "search_facebook_users")
    t_get     = next(t for t in tools if t.name == "get_facebook_profile")

    rl = await t_search.ainvoke({"q": name, "limit": 10, "fuzzy": True})
    users = _tool_to_obj(rl)
    if not users:
        return None

    # For Facebook, prioritize exact matches because years of experience are unavailable.
    best = sorted(users, key=lambda u: (u.get("match_type") != "exact"))[0]
    prof_raw = await t_get.ainvoke({"user_id": best["id"]})
    profile = _tool_to_obj(prof_raw)
    return {"candidate": best, "profile": profile}

import re
import json
from langchain_core.messages import HumanMessage

def extract_year_hints(cv_text: str, max_hints: int = 12) -> list[str]:
    t = cv_text or ""
    # capture "YYYY - YYYY" or "YYYY - Present"
    pattern = re.compile(r"\b((?:(19|20)\d{2})\s*[-]\s*(?:(19|20)\d{2})|present)\b", re.IGNORECASE)
    range_strs = []
    for m in pattern.finditer(t):
        a = m.group(1)
        b = m.group(2)
        range_strs.append(f"(a)-(b)")

    # also capture single years (YYYY)
    years = re.findall(r"\b(?:(19|20)\d{2})\b", t)

    # combine + dedupe
    seen = set()
    out = []
    for year in years:
        if year not in seen:
            out.append(year)
            seen.add(year)
    return out[:max_hints]

```

```

for x in range_strs + years:
    x = x.strip()
    if x and x not in seen:
        seen.add(x)
        out.append(x)

return out[:max_hints]

def extract_cv_json(llm, cv_text: str) -> dict:
    year_hints = extract_year_hints(cv_text)

    prompt = f"""
You are an information extraction system. Return STRICT JSON only (no markdown).
Your job is to extract structured fields from a CV.

IMPORTANT RULES:
1) For every experience item, include start_year and end_year if stated. If not stated, use null.
2) If the CV says "Present", set end_year to null and is_current=true.
3) Keep company/title/school strings as they appear (do not invent).

Return JSON with this exact schema:
{{
    "name": "",
    "city": "",
    "country": "",
    "education": [{"school": "", "degree": "", "field": "", "start_year": null, "end_year": null}],
    "experience": [{"company": "", "title": "", "start_year": null, "end_year": null, "is_current": false}],
    "skills": []
}}
"""

Year hints extracted from the CV (may help you locate dates): {year_hints}

CV TEXT:
{cv_text}
"""

resp = llm.invoke([HumanMessage(content=prompt)])
text = resp.content
start = text.find("(")
end = text.rfind(")")
data = json.loads(text[start:end+1])
return data

import pandas as pd
import re

def summarize_cv(cv: dict) -> dict:
    exp = cv.get("experience") or []
    edu = cv.get("education") or []
    # count years coverage
    exp_with_years = sum(1 for e in exp if isinstance(e.get("start_year"), int) or isinstance(e.get("end_year"), int))
    edu_with_years = sum(1 for e in edu if isinstance(e.get("start_year"), int) or isinstance(e.get("end_year"), int))

    return {
        "name": cv.get("name"),
        "city": cv.get("city"),
        "country": cv.get("country"),
        "n_exp": len(exp),
        "n_exp_with_years": exp_with_years,
        "n_edu": len(edu),
        "n_edu_with_years": edu_with_years,
        "n_skills": len(cv.get("skills") or [])
    }

def audit_extraction(all_csvs, llm):
    rows = []
    cv_objects = []
    for item in all_csvs:
        cv = extract_cv_json(llm, item["text"])
        cv_objects.append(cv)
        row = {"file": item["file"], **summarize_cv(cv)}
        rows.append(row)
    df = pd.DataFrame(rows)
    return df, cv_objects

df_extract, extracted_csvs = audit_extraction(all_csvs, llm)
df_extract
def normalize_cv(cv: dict, raw_text: str) -> dict:
    """
    Postprocess to improve year coverage:
    - If experience has company/title but missing years, try to find a nearby year range in raw_text.
    This is heuristic but helps avoid CV^1 artifacts.
    """
    t = raw_text.replace("–", "-").replace("—", "-")
    year_range_pat = re.compile(r"\b((?:(19|20)\d{2})\s*[-]\s*(?:(19|20)\d{2})|present)\b", re.IGNORECASE)

```

```

# If many experience items missing years, try a coarse fallback:
ranges = year_range_pat.findall(t)
# ranges are tuples; reconstruct
range_strs = re.findall(r"\b(?:19|20)\d{2}\s*\-\s*(?:19|20)\d{2}\b|present\b", t, flags=re.IGNORECASE)
# easier: directly capture as strings
range_strs = re.findall(r"\b(?:19|20)\d{2}\b|\s*\-\s*(?:19|20)\d{2}\b|present\b", t, flags=re.IGNORECASE)

exp = cv.get("experience") or []
missing = [i for i, e in enumerate(exp) if (e.get("company") or e.get("title")) and e.get("start_year") is None]
if missing and range_strs:
    # assign ranges in order to missing slots (best-effort)
    for idx, (sy, ey) in zip(missing, range_strs):
        try:
            exp[idx]["start_year"] = int(sy)
        except:
            pass
        if str(ey).lower() == "present":
            exp[idx]["end_year"] = None
            exp[idx]["is_current"] = True
        else:
            try:
                exp[idx]["end_year"] = int(ey)
            except:
                pass
cv["experience"] = exp
return cv

def internal_inconsistency_penalty(cv: dict) -> tuple[float, list[str]]:
    """
    TA requirement: internal contradictions should only result in a score deduction and must not trigger an early rejection.
    Here we apply a light penalty (up to 0.15) to avoid mistakenly rejecting a genuine CV.
    """
    issues = []
    penalty = 0.0

    # Time overlap: if multiple roles are marked as current or the dates clearly conflict.
    exps = cv.get("experience", []) or []
    current_count = sum(1 for e in exps if e.get("is_current") is True)
    if current_count >= 2:
        issues.append("Multiple current jobs in CV (possible timeline inconsistency).")
        penalty += 0.08

    # Location: city/country is missing, but multiple place names are detected.
    if (not cv.get("city") or not cv.get("country")) and len(re.findall(r"\b(Hong Kong|Singapore|Tokyo|London|Kowloon|Philippines)", cv["location"])) > 1:
        issues.append("Ambiguous/mixed location signals in CV (weak internal inconsistency).")
        penalty += 0.04

    # maximum
    penalty = min(penalty, 0.15)
    return penalty, issues

def compute_score(cv: dict, li: dict | None, fb: dict | None) -> tuple[float, list[str]]:
    import json

    issues = []
    major = 0
    minor = 0

    def _soft_mismatch(a, b):
        a = (a or "").lower().strip()
        b = (b or "").lower().strip()
        if not a or not b:
            return False
        return (a not in b) and (b not in a)

    def _latest_role(exps):
        exps = exps or []
        exps_sorted = sorted(
            exps,
            key=lambda e: (e.get("start_year") is None, -(e.get("start_year") or 0)))
        ) return exps_sorted[0] if exps_sorted else None

    # =====
    # LinkedIn (primary verification)
    # =====

    li_prof = None
    industry_conflict = False

    if not li:
        issues.append("LinkedIn profile not found after fallback search.")
        minor += 1
    else:
        li_prof = li["profile"]
        match_type = (li.get("candidate") or {}).get("match_type")
        is_exact = (match_type == "exact")

        cv_exps = cv.get("experience") or []
        li_exps = li_prof.get("experience") or []

        cv_role = next((e for e in cv_exps if e.get("is_current")), None) or _latest_role(cv_exps)
        li_role = next((e for e in li_exps if e.get("is_current")), None) or _latest_role(li_exps)

        company_mis = False
        title_mis = False

        if cv_role and li_role:
            if _soft_mismatch(cv_role.get("company"), li_role.get("company")):
                issues.append(f"Current/Latest company mismatch: CV={cv_role.get('company')} vs LI={li_role.get('company')}")
            company_mis = True

            if _soft_mismatch(cv_role.get("title"), li_role.get("title")):
                issues.append(f"Current/Latest title mismatch: CV={cv_role.get('title')} vs LI={li_role.get('title')}")
            title_mis = True

        # Title mismatch is weak signal
        if title_mis:
            minor += 1
            issues.append("Treated title mismatch as minor (title wording/noise or injected inconsistency).")

        # Education mismatch (weak)
        cv_schools = [
            e.get("school", "").lower().strip()
            for e in (cv.get("education") or [])
            if e.get("school")
        ]
        li_schools = [
            e.get("school", "").lower().strip()
            for e in (li_prof.get("education") or [])
            if e.get("school")
        ]
        edu_mis = bool(cv_schools and li_schools and cv_schools.isdisjoint(li_schools))

        if edu_mis:
            issues.append("Education schools mismatch between CV and LinkedIn.")
            if is_exact:
                minor += 1
            else:
                issues.append("Education mismatch not penalized because LinkedIn match is not exact.")

    # -----
    # Robust industry/domain mismatch
    # -----

    li_text = f"{li_prof.get('headline', '')} {li_prof.get('industry', '')}.lower()"
    cv_text = json.dumps(cv).lower()

    # domain keyword sets
    legal_kw = ["legal", "law", "litigation", "compliance", "contract", "paralegal"]
    tech_kw = ["engineer", "software", "developer", "data", "scientist", "machine learning", "ai", "cloud", "quantum"]
    mkt_kw = ["marketing", "seo", "social media", "content", "brand", "growth"]
    ops_kw = ["logistics", "supply chain", "operations", "ops", "warehouse", "procurement"]

    li_is_legal = any(k in li_text for k in legal_kw)
    li_is_tech = any(k in li_text for k in tech_kw)
    li_is_ops = any(k in li_text for k in ops_kw)
    li_is_mkt = any(k in li_text for k in mkt_kw)

    cv_is_legal = any(k in cv_text for k in legal_kw)
    cv_is_tech = any(k in cv_text for k in tech_kw)
    cv_is_ops = any(k in cv_text for k in ops_kw)
    cv_is_mkt = any(k in cv_text for k in mkt_kw)

    industry_conflict = False

    # Strong conflicts -> MAJOR
    if (li_is_legal and cv_is_tech) or (li_is_tech and cv_is_legal):
        issues.append("Strong conflicts between LinkedIn and CV industries: Legal vs Tech")
        industry_conflict = True

```

```

issues.append("LinkedIn profile found after fallback search.")
minor += 1

else:
    li_prof = li["profile"]
    match_type = (li.get("candidate") or {}).get("match_type")
    is_exact = (match_type == "exact")

    cv_exps = cv.get("experience") or []
    li_exps = li_prof.get("experience") or []

    cv_role = next((e for e in cv_exps if e.get("is_current")), None) or _latest_role(cv_exps)
    li_role = next((e for e in li_exps if e.get("is_current")), None) or _latest_role(li_exps)

    company_mis = False
    title_mis = False

    if cv_role and li_role:
        if _soft_mismatch(cv_role.get("company"), li_role.get("company")):
            issues.append(f"Current/Latest company mismatch: CV={cv_role.get('company')} vs LI={li_role.get('company')}")
        company_mis = True

        if _soft_mismatch(cv_role.get("title"), li_role.get("title")):
            issues.append(f"Current/Latest title mismatch: CV={cv_role.get('title')} vs LI={li_role.get('title')}")
        title_mis = True

    # Title mismatch is weak signal
    if title_mis:
        minor += 1
        issues.append("Treated title mismatch as minor (title wording/noise or injected inconsistency).")

    # Education mismatch (weak)
    cv_schools = [
        e.get("school", "").lower().strip()
        for e in (cv.get("education") or [])
        if e.get("school")
    ]
    li_schools = [
        e.get("school", "").lower().strip()
        for e in (li_prof.get("education") or [])
        if e.get("school")
    ]
    edu_mis = bool(cv_schools and li_schools and cv_schools.isdisjoint(li_schools))

    if edu_mis:
        issues.append("Education schools mismatch between CV and LinkedIn.")
        if is_exact:
            minor += 1
        else:
            issues.append("Education mismatch not penalized because LinkedIn match is not exact.")

    # -----
    # Robust industry/domain mismatch
    # -----

    li_text = f"{li_prof.get('headline', '')} {li_prof.get('industry', '')}.lower()"
    cv_text = json.dumps(cv).lower()

    # domain keyword sets
    legal_kw = ["legal", "law", "litigation", "compliance", "contract", "paralegal"]
    tech_kw = ["engineer", "software", "developer", "data", "scientist", "machine learning", "ai", "cloud", "quantum"]
    mkt_kw = ["marketing", "seo", "social media", "content", "brand", "growth"]
    ops_kw = ["logistics", "supply chain", "operations", "ops", "warehouse", "procurement"]

    li_is_legal = any(k in li_text for k in legal_kw)
    li_is_tech = any(k in li_text for k in tech_kw)
    li_is_ops = any(k in li_text for k in ops_kw)
    li_is_mkt = any(k in li_text for k in mkt_kw)

    cv_is_legal = any(k in cv_text for k in legal_kw)
    cv_is_tech = any(k in cv_text for k in tech_kw)
    cv_is_ops = any(k in cv_text for k in ops_kw)
    cv_is_mkt = any(k in cv_text for k in mkt_kw)

    industry_conflict = False

    # Strong conflicts -> MAJOR
    if (li_is_legal and cv_is_tech) or (li_is_tech and cv_is_legal):
        issues.append("Strong conflicts between LinkedIn and CV industries: Legal vs Tech")
        industry_conflict = True

```

```

issues.append("Strong domain conflict: Legal vs Tech between CV and LinkedIn.")
if is_exact:
    industry_conflict = True
    major += 1
else:
    minor += 1
    issues.append("Domain conflict treated as minor because LinkedIn match is not exact.")
# Weak conflicts -> MINOR (do not reject)
elif (li_is_ops and cv_is_mkt) or (li_is_mkt and cv_is_ops):
    issues.append("Weak domain conflict: Marketing vs Logistics/Operations between CV and LinkedIn.")
    minor += 1

# -----
# Company mismatch logic (conditional)
# -----

if company_mis:
    if industry_conflict:
        issues.append(
            "Escalation: company mismatch supported by domain conflict."
        )
        major += 1
    else:
        minor += 1
        issues.append(
            "Treated company mismatch as minor (possible injected current-job inconsistency)."
        )

# -----
# Facebook (very weak auxiliary)
# -----

if fb:
    fb_prof = fb["profile"]
    cv_city = (cv.get("city") or "").lower().strip()
    fb_city = (fb_prof.get("city") or "").lower().strip()

    if cv_city and fb_city and (cv_city not in fb_city) and (fb_city not in cv_city):
        issues.append(
            f"(Weak) City mismatch: CV={cv.get('city')} vs FB={fb_prof.get('city')}"
        )

# -----
# Internal inconsistency penalty
# -----

pen, pen_issues = internal_inconsistency_penalty(cv)
issues.extend(pen_issues)

# -----
# Final scoring
# -----

raw = 0.85 - 0.40 * major - 0.08 * minor - pen
score = max(0.05, min(1.0, raw))

return float(score), issues

```

```

import asyncio

async def run_all(all_csvs, tools, llm):
    scores = []
    debug = []

    for item in all_csvs:
        cv = extract_cv_json(llm, item["text"])
        cv = normalize_cv(cv, item["text"])
        name = (cv.get("name") or "").strip()

        # LinkedIn location hint: 可选, 有就用
        loc = None
        if cv.get("city") and cv.get("country"):
            loc = f'{cv["city"]}, {cv["country"]}'
        elif cv.get("country"):
            loc = cv["country"]

        li = await fetch_linkedin_best(tools, name, loc) if name else None
        fb = await fetch_facebook_best(tools, name) if name else None

        score, issues = compute_score(cv, li, fb)

```

```

scores.append(score)
debug.append({
    "file": item["file"],
    "name": name,
    "score": score,
    "issues": issues,
    "li_match_type": (li.get("candidate") or {}).get("match_type") if li else None,
    "li_headline": (li.get("profile") or {}).get("headline") if li else None,
    "li_industry": (li.get("profile") or {}).get("industry") if li else None,
})
}

return scores, debug

scores, debug = await run_all(all_csvs, mcp_tools, 11m)
scores

```

[0.53, 0.61, 0.61, 0.05, 0.05]

Evaluation code

In the test phase, you will be given 5 CV files with fixed names:

CV_1.pdf, CV_2.pdf, CV_3.pdf, CV_4.pdf, CV_5.pdf

Your system must process these CVs and output a list of 5 scores, one score per CV, in the same order:

scores = [s1, s2, s3, s4, s5]

Each score must be a float in the range [0, 1], representing the reliability or confidence that the CV is valid (or meets the task criteria).

The ground-truth labels are binary:

groundtruth = [0 or 1, ..., 0 or 1]

Each CV is evaluated independently using a threshold of 0.5:

- If score > 0.5 and groundtruth == 1 → Full credit
- If score ≤ 0.5 and groundtruth == 0 → Full credit
- Otherwise → No credit

In other words, 0.5 is the decision threshold.

- Each CV contributes equally.
- Final score = (number of correct decisions) / 5

```

# -----
# Evaluation code
# -----

def evaluate(scores, groundtruth, threshold=0.5):
    """
    scores: list of floats in [0, 1], length = 5
    groundtruth: list of ints (0 or 1), length = 5
    """

    assert len(scores) == 5
    assert len(groundtruth) == 5

    correct = 0
    decisions = []

    for s, gt in zip(scores, groundtruth):
        pred = 1 if s > threshold else 0
        decisions.append(pred)
        if pred == gt:
            correct += 1

    final_score = correct / len(scores)

    return {
        "decisions": decisions,
        "correct": correct,
        "total": len(scores),
        "final_score": final_score
    }
}

```

```
groundtruth = [1, 1, 1, 0, 0] # Do not modify  
result = evaluate(scores, groundtruth)  
print(result)  
  
{'decisions': [1, 1, 1, 0, 0], 'correct': 5, 'total': 5, 'final_score': 1.0}
```