

```
# =====
# 1. Install dependencies
# =====
!pip install -U openai numpy pandas scipy tabulate matplotlib

Requirement already satisfied: openai in /usr/local/lib/python3.12/dist-packages (2.24.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (2.4.2)
Requirement already satisfied: pandas in /usr/local/lib/python3.12/dist-packages (3.0.1)
Requirement already satisfied: scipy in /usr/local/lib/python3.12/dist-packages (1.17.1)
Requirement already satisfied: tabulate in /usr/local/lib/python3.12/dist-packages (0.9.0)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages (3.10.8)
Requirement already satisfied: aiohttp<5,>=3.5.0 in /usr/local/lib/python3.12/dist-packages (from openai) (4.12.1)
Requirement already satisfied: distro<2,>=1.7.0 in /usr/local/lib/python3.12/dist-packages (from openai) (1.9.0)
Requirement already satisfied: httpx<1,>=0.23.0 in /usr/local/lib/python3.12/dist-packages (from openai) (0.28.1)
Requirement already satisfied: jiter<1,>=0.10.0 in /usr/local/lib/python3.12/dist-packages (from openai) (0.13.0)
Requirement already satisfied: pydantic<3,>=1.9.0 in /usr/local/lib/python3.12/dist-packages (from openai) (2.12.3)
Requirement already satisfied: sniffio in /usr/local/lib/python3.12/dist-packages (from openai) (1.3.1)
Requirement already satisfied: tqdm>4 in /usr/local/lib/python3.12/dist-packages (from openai) (4.67.3)
Requirement already satisfied: typing_extensions<5,>=4.11 in /usr/local/lib/python3.12/dist-packages (from openai) (4.15.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.3.3)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (4.61.1)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.4.9)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (26.0)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (11.3.0)
Requirement already satisfied: pyparsing>=3 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (3.3.2)
Requirement already satisfied: idna>=2.8 in /usr/local/lib/python3.12/dist-packages (from aiohttp<5,>=3.5.0->openai) (3.11)
Requirement already satisfied: certifi in /usr/local/lib/python3.12/dist-packages (from httpx<1,>=0.23.0->openai) (2026.1.4)
Requirement already satisfied: httpcore==1.* in /usr/local/lib/python3.12/dist-packages (from httpx<1,>=0.23.0->openai) (1.0.9)
Requirement already satisfied: h11>=0.16 in /usr/local/lib/python3.12/dist-packages (from httpcore==1.*->httpx<1,>=0.23.0->openai) (0.16)
Requirement already satisfied: annotated-types>=0.6.0 in /usr/local/lib/python3.12/dist-packages (from pydantic<3,>=1.9.0->openai) (0.7)
Requirement already satisfied: pydantic-core==2.41.4 in /usr/local/lib/python3.12/dist-packages (from pydantic<3,>=1.9.0->openai) (2.41)
Requirement already satisfied: typing-inspection>=0.4.2 in /usr/local/lib/python3.12/dist-packages (from pydantic<3,>=1.9.0->openai) (0.7)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2->pandas) (1.17.0)
```

```
# =====
# 2. Import library
# =====
import os
import time
import json
import math
import random
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from openai import OpenAI
from scipy.stats import norm
from tabulate import tabulate
```

```
# =====
# 3. Configure DeepSeek API
# =====
from openai import OpenAI
from google.colab import userdata
import random
import numpy as np

DEEPESEEK_API_KEY = userdata.get("DEEPESEEK_API_KEY")

if not DEEPESEEK_API_KEY:
    raise ValueError("DEEPESEEK_API_KEY is empty. Please set it in Colab Secrets first.")

client = OpenAI(
    api_key=DEEPESEEK_API_KEY,
    base_url="https://api.deepseek.com"
)

MODEL_NAME = "deepseek-chat"

# Baseline: 1
# Modification: 0
TEMPERATURE = 0

# To test whether it can run, it is recommended to start with a small scale.
N_LIST = [1, 3]
BLOCKS = 2
```

```
# common settings for verbal tasks in the original paper
ALPHABET = "bcdfghjklnpqrstuvwxyz"
SEQ_LENGTH = 24
MATCHES = 8

# Fixed random seed for easy reproduction
RANDOM_SEED = 42
random.seed(RANDOM_SEED)
np.random.seed(RANDOM_SEED)

# Output filename
OUTPUT_FILE = f"all_trials_verbal_temp{TEMPERATURE}.json"

print("DeepSeek client ready")
print("model:", MODEL_NAME)
print("temperature:", TEMPERATURE)
print("blocks:", BLOCKS)
print("key length:", len(DEEPSEEK_API_KEY))
print("key prefix:", DEEPSEEK_API_KEY[:5])
```

```
DeepSeek client ready
model: deepseek-chat
temperature: 0
blocks: 2
key length: 35
key prefix: sk-bb
```

```
# =====
# 4. Generate verbal n-back sequences
# =====

def generate_nback_sequences(n, alphabet, seq_length, matches, num_sequences):
    sequences = []

    for _ in range(num_sequences):
        seq = []
        conditions = []

        match_positions = random.sample(range(n, seq_length), matches)
        match_positions.sort()

        for i in range(seq_length):
            if i in match_positions:
                # Forced construction of match
                seq.append(seq[i - n])
                conditions.append("m")
            else:
                # Construct a non-match to avoid accidental n-backmatches
                random_letter = random.choice(alphabet)
                while i >= n and random_letter == seq[i - n]:
                    random_letter = random.choice(alphabet)

                seq.append(random_letter)
                conditions.append("-")

        sequences.append(("".join(seq), "".join(conditions)))

    return sequences

def build_all_trials(n_list, blocks, alphabet, seq_length, matches):
    all_trials = {}

    for n in n_list:
        sequences = generate_nback_sequences(
            n=n,
            alphabet=alphabet,
            seq_length=seq_length,
            matches=matches,
            num_sequences=blocks
        )

        for b, (seq, cond) in enumerate(sequences):
            trials = []
            for i in range(len(seq)):
                trials.append({
                    "stimulus": seq[i],
                    "target": cond[i],
                    "response": "",
                    "correct": None,
                    "rt": None
                })

            all_trials[f"{n}back_{b}"] = trials
```

```

        return all_trials

all_trials = build_all_trials(
    n_list=N_LIST,
    blocks=BLOCKS,
    alphabet=ALPHABET,
    seq_length=SEQ_LENGTH,
    matches=MATCHES
)

print("Data preparation complete.")
print("keys Example: ", list(all_trials.keys())[:3])

Data preparation complete.
keys Example: ['1back_0', '1back_1', '3back_0']

# =====
# 5. DeepSeek function calls (stable version)
# =====
def normalize_response(text):
    """
    Try to normalize the model output to 'm' or '-'.

    if text is None:
        return None

    text = str(text).strip().lower()

    # Exact match
    if text == "m":
        return "m"
    if text == "-":
        return "-"

    # Common multi-output scenarios
    text = text.replace(" ", "")
    if text.startswith("m"):
        return "m"
    if text.startswith("-"):
        return "-"

    # Iterate through the characters to find the first valid character.
    for ch in text:
        if ch in ["m", "-"]:
            return ch

    return None

def get_model_response(messages, model_name, temperature, max_retries=3, sleep_seconds=2):
    if "client" not in globals():
        raise ValueError("The client has not been initialized yet. Please run the API configuration section first")

    last_error = None

    for attempt in range(max_retries):
        try:
            start_time = time.time()

            completion = client.chat.completions.create(
                model=model_name,
                messages=messages,
                temperature=temperature,
                max_tokens=5
            )

            response_time = time.time() - start_time
            raw_text = completion.choices[0].message.content
            response = normalize_response(raw_text)

            return response, raw_text, response_time

        except Exception as e:
            last_error = e
            print(f"[API ERROR] attempt {attempt+1}/{max_retries}: {e}")
            time.sleep(sleep_seconds)

    raise RuntimeError(f"The API kept failing, eventually resulting in an error.: {last_error}")

```

```

# =====
# 6. Run a single n-back task (stable version)
# =====

def run_nback_task(all_trials, n, blocks, model_name, temperature, output_file, sleep_between_blocks=1):
    if n == 1:
        task_instruction = (
            "You are asked to perform a 1-back task. "
            "You will see a sequence of letters one by one. "
            "Respond with 'm' whenever the current letter is the same as the previous letter, "
            "and '-' otherwise. Only output 'm' or '-'. No explanation."
        )
    elif n == 2:
        task_instruction = (
            "You are asked to perform a 2-back task. "
            "You will see a sequence of letters one by one. "
            "Respond with 'm' whenever the current letter is the same as the letter two steps ago, "
            "and '-' otherwise. Only output 'm' or '-'. No explanation."
        )
    elif n == 3:
        task_instruction = (
            "You are asked to perform a 3-back task. "
            "You will see a sequence of letters one by one. "
            "Respond with 'm' whenever the current letter is the same as the letter three steps ago, "
            "and '-' otherwise. Only output 'm' or '-'. No explanation."
        )
    else:
        raise ValueError("Currently only supports n=1/2/3")

    for b in range(blocks):
        trial_key = f"{n}back_{b}"
        messages = [{"role": "system", "content": task_instruction}]

        print(f"\n===== Running {trial_key} =====")

        for i, trial in enumerate(all_trials[trial_key]):
            stimulus = trial["stimulus"]
            target = trial["target"]

            print(f"block {b}, trial {i}: stimulus={stimulus}, target={target}")

            messages.append({"role": "user", "content": stimulus})

            try:
                response, raw_text, rt = get_model_response(
                    messages=messages,
                    model_name=model_name,
                    temperature=temperature
                )

                trial["rt"] = rt
                trial["raw_output"] = raw_text

                # If the parsing still fails, provide a default value to avoid interrupting the entire experiment
                if response not in ["m", "-"]:
                    print(f"[WARNING] Unable to recognize output, raw={raw_text}, Default is recorded as '-'")
                    response = "-"

                trial["response"] = response
                trial["correct"] = (response == target)

                print(f"model raw: {raw_text}")
                print(f"parsed: {response}")
                print(f"rt: {rt:.2f}s")
                print(f"correct" if trial["correct"] else "incorrect")
                print("-----")

                messages.append({"role": "assistant", "content": response})
            except Exception as e:
                print(f"[TRIAL ERROR] block {b}, trial {i}: {e}")

                trial["rt"] = None
                trial["raw_output"] = str(e)
                trial["response"] = "ERROR"
                trial["correct"] = False

                # To avoid interrupting the entire experiment, a default answer is added to the context.
                messages.append({"role": "assistant", "content": "-"})


        # Save the data after each block is completed to prevent data loss if the process is interrupted.
        with open(output_file, "w", encoding="utf-8") as f:
            json.dump(all_trials, f, ensure_ascii=False, indent=2)

```

```

        print(f"[AUTO-SAVED] {output_file}")
        time.sleep(sleep_between_blocks)

    return all_trials

# =====
# 7. Officially running experiment (stable version)
# =====
print("Start running the experiment...")
print(f"MODEL_NAME = {MODEL_NAME}")
print(f"TEMPERATURE = {TEMPERATURE}")
print(f"N_LIST = {N_LIST}")
print(f"BLOCKS = {BLOCKS}")
print(f"OUTPUT_FILE = {OUTPUT_FILE}")

for n in N_LIST:
    all_trials = run_nback_task(
        all_trials=all_trials,
        n=n,
        blocks=BLOCKS,
        model_name=MODEL_NAME,
        temperature=TEMPERATURE,
        output_file=OUTPUT_FILE,
        sleep_between_blocks=1
    )

with open(OUTPUT_FILE, "w", encoding="utf-8") as f:
    json.dump(all_trials, f, ensure_ascii=False, indent=2)

print(f"\nThe experiment is complete, and the results have been saved to: {OUTPUT_FILE}")

Start running the experiment...
MODEL_NAME = deepseek-chat
TEMPERATURE = 0
N_LIST = [1, 3]
BLOCKS = 2
OUTPUT_FILE = all_trials_verbal_temp0.json

===== Running lback_0 =====
block 0, trial 0: stimulus=x, target=-
model raw: -
parsed: -
rt: 2.27s
correct

-----
block 0, trial 1: stimulus=x, target=m
model raw: -
parsed: -
rt: 1.87s
incorrect

-----
block 0, trial 2: stimulus=d, target=-
model raw: -
parsed: -
rt: 1.68s
correct

-----
block 0, trial 3: stimulus=y, target=-
model raw: -
parsed: -
rt: 1.62s
correct

-----
block 0, trial 4: stimulus=y, target=m
model raw: m
parsed: m
rt: 1.22s
correct

-----
block 0, trial 5: stimulus=y, target=m
model raw: m
parsed: m
rt: 1.58s
correct

-----
block 0, trial 6: stimulus=s, target=-
model raw: -
parsed: -
rt: 1.15s
correct

-----
block 0, trial 7: stimulus=c, target=-
model raw: -
parsed: -
rt: 1.34s

```

correct

block 0, trial 8: stimulus=c, target=m

```

# =====
# 8. Calculate summary stats
# =====
def compute_summary_stats(all_trials, n_list, blocks):
    summary_stats = {}
    raw_data = {}

    for n in n_list:
        hit_rate = []
        false_alarm_rate = []
        accuracy = []
        response_time = []
        d_prime = []

        for b in range(blocks):
            trials = all_trials[f"{n}back_{b}"]

            hits = 0
            false_alarms = 0
            total_targets = 0
            total_lures = 0
            total_rt = 0

            for trial in trials:
                if trial["target"] == "m":
                    total_targets += 1
                    if trial["correct"]:
                        hits += 1
                        total_rt += trial["rt"]
                else:
                    total_lures += 1
                    if trial["response"] == "m":
                        false_alarms += 1

            hr = (hits / total_targets) * 100 if total_targets > 0 else np.nan
            far = (false_alarms / total_lures) * 100 if total_lures > 0 else np.nan
            acc = ((hits + (total_lures - false_alarms)) / (total_targets + total_lures)) * 100

            # d'
            hit_rate_adjusted = np.clip(hits / total_targets, 0.01, 0.99) if total_targets > 0 else np.nan
            false_alarm_rate_adjusted = np.clip(false_alarms / total_lures, 0.01, 0.99) if total_lures > 0 else np.nan

            if not np.isnan(hit_rate_adjusted) and not np.isnan(false_alarm_rate_adjusted):
                dp = abs(norm.ppf(hit_rate_adjusted) - norm.ppf(false_alarm_rate_adjusted))
            else:
                dp = np.nan

            mean_rt = total_rt / hits if hits > 0 else np.nan

            hit_rate.append(hr)
            false_alarm_rate.append(far)
            accuracy.append(acc)
            response_time.append(mean_rt)
            d_prime.append(dp)

        raw_data[n] = {
            "hit_rate": hit_rate,
            "false_alarm_rate": false_alarm_rate,
            "accuracy": accuracy,
            "response_time": response_time,
            "d_prime": d_prime
        }

        summary_stats[n] = {
            "hit_rate": {
                "mean": np.nanmean(hit_rate),
                "stderr": np.nanstd(hit_rate, ddof=1) / math.sqrt(blocks) if blocks > 1 else 0
            },
            "false_alarm_rate": {
                "mean": np.nanmean(false_alarm_rate),
                "stderr": np.nanstd(false_alarm_rate, ddof=1) / math.sqrt(blocks) if blocks > 1 else 0
            },
            "accuracy": {
                "mean": np.nanmean(accuracy),
                "stderr": np.nanstd(accuracy, ddof=1) / math.sqrt(blocks) if blocks > 1 else 0
            },
            "d_prime": {
                "mean": np.nanmean(d_prime),
                "stderr": np.nanstd(d_prime, ddof=1) / math.sqrt(blocks) if blocks > 1 else 0
            }
        }
    }

```

```

        "stderr": np.nanstd(d_prime, ddof=1) / math.sqrt(blocks) if blocks > 1 else 0
    }
}

return summary_stats, raw_data

```

def create\_table(summary\_stats, n\_list):
 headers = ["N-back", "Hit Rate (%)", "False Alarm Rate (%)", "Accuracy (%)", "D Prime"]
 rows = []
 for n in n\_list:
 rows.append([
 f'{n}-back',
 f'{summary\_stats[n]["hit\_rate"]["mean"]:.2f} ± {summary\_stats[n]["hit\_rate"]["stderr"]:.2f}',
 f'{summary\_stats[n]["false\_alarm\_rate"]["mean"]:.2f} ± {summary\_stats[n]["false\_alarm\_rate"]["stderr"]:.2f}',
 f'{summary\_stats[n]["accuracy"]["mean"]:.2f} ± {summary\_stats[n]["accuracy"]["stderr"]:.2f}',
 f'{summary\_stats[n]["d\_prime"]["mean"]:.2f} ± {summary\_stats[n]["d\_prime"]["stderr"]:.2f}',
 ])
 print(tabulate(rows, headers=headers, tablefmt="github"))

 df = pd.DataFrame(rows, columns=headers)
 return df

```

summary_stats, raw_data = compute_summary_stats(all_trials, N_LIST, BLOCKS)
result_df = create_table(summary_stats, N_LIST)
result_df

```

N-back	Hit Rate (%)	False Alarm Rate (%)	Accuracy (%)	D Prime
1-back	87.50 ± 0.00	0.00 ± 0.00	95.83 ± 0.00	3.48 ± 0.00
3-back	25.00 ± 25.00	0.00 ± 0.00	75.00 ± 8.33	1.16 ± 1.16

  

	N-back	Hit Rate (%)	False Alarm Rate (%)	Accuracy (%)	D Prime
0	1-back	87.50 ± 0.00	0.00 ± 0.00	95.83 ± 0.00	3.48 ± 0.00
1	3-back	25.00 ± 25.00	0.00 ± 0.00	75.00 ± 8.33	1.16 ± 1.16

```

# =====
# 9. Visualization accuracy
# =====
x = [str(n) for n in N_LIST]
y = [summary_stats[n]["accuracy"]["mean"] for n in N_LIST]
yerr = [summary_stats[n]["accuracy"]["stderr"] for n in N_LIST]

plt.figure(figsize=(6, 4))
plt.bar(x, y, yerr=yerr)
plt.xlabel("N-back")
plt.ylabel("Accuracy (%)")
plt.title(f"Verbal N-Back Accuracy ({MODEL_NAME}, temp={TEMPERATURE})")
plt.show()

```



